

// ContainerDay Workshops 2023 - Munich

# Mastering Production Grade Deployments

K8s Deep Dive



November 23 | Christian Hörl

# About



## Estd. 2007

- based in Berlin
- 3 Datacenter regions in Germany
- Managed Hosting, Openstack
- Managed Kubernetes with "MetaKube"



# Topics

- **K8s on-board features**

- Pod Design: Good Practices
- Deployment Strategy & Affinity
- Resources: Requests, Limits, Priority
- Scaling
- Disruption Budget, Security

- **3rd-party Tools**

- external-dns, cert-manager
- ingress-nginx & canary deployments
- Helm & Helmfile

Production-grade  
Deployments?

**Production readiness?**

# Redundancy

# Redundancy

- having “more than 1”
- on different layers
- K8s *Nodes* and *Replicas*

# High Load & Scaling



# High Load & Scaling

- react
- scale fast
- automate
- *K8s Pod / Node auto-scaling*

# Persistence

# Persistence

- centralized, decentralized
- distribute persistent data
- move with the consumer
- K8s *persistentVolumes*

# Self-healing Architecture

# Security

# Security

- enable, configure, use
- intercept workload behavior
- K8s *securityContext*

# Backup & Recovery

# Quality Assurance



# The Pillars of Production-Readiness

K8s features



3rd-party tools



reactive &  
supportive



automation & continuous deployment



# K8s features



## Get ready



**Hands-on repo**

[github.com/syseleven/containerday-workshops](https://github.com/syseleven/containerday-workshops)



**Kubeconfig**

[sys11.it/production-grade-2023](https://sys11.it/production-grade-2023)

# Pod Configuration Best Practices


# Deployment Strategy



# Deployment strategy: Recreate



# Deployment strategy: Rolling Update




# Hands-on


<https://github.com/syseleven/containerday-workshops/tree/main/100-deployment-strategy>



# Affinity





# Node Affinity




controlled by Node labels

# Pod Affinity



controlled by Pod labels

# Pod AntiAffinity




controlled by Pod labels and topology key

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/105-affinity>

# Resource Requests & Limits



# Hands-on


<https://github.com/syseleven/containerday-workshops/tree/main/110-resource-requests-limits>

# PriorityClasses



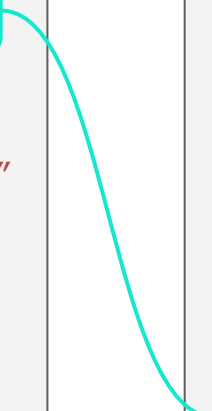


## PriorityClass - Faster scheduling of pods with higher priority



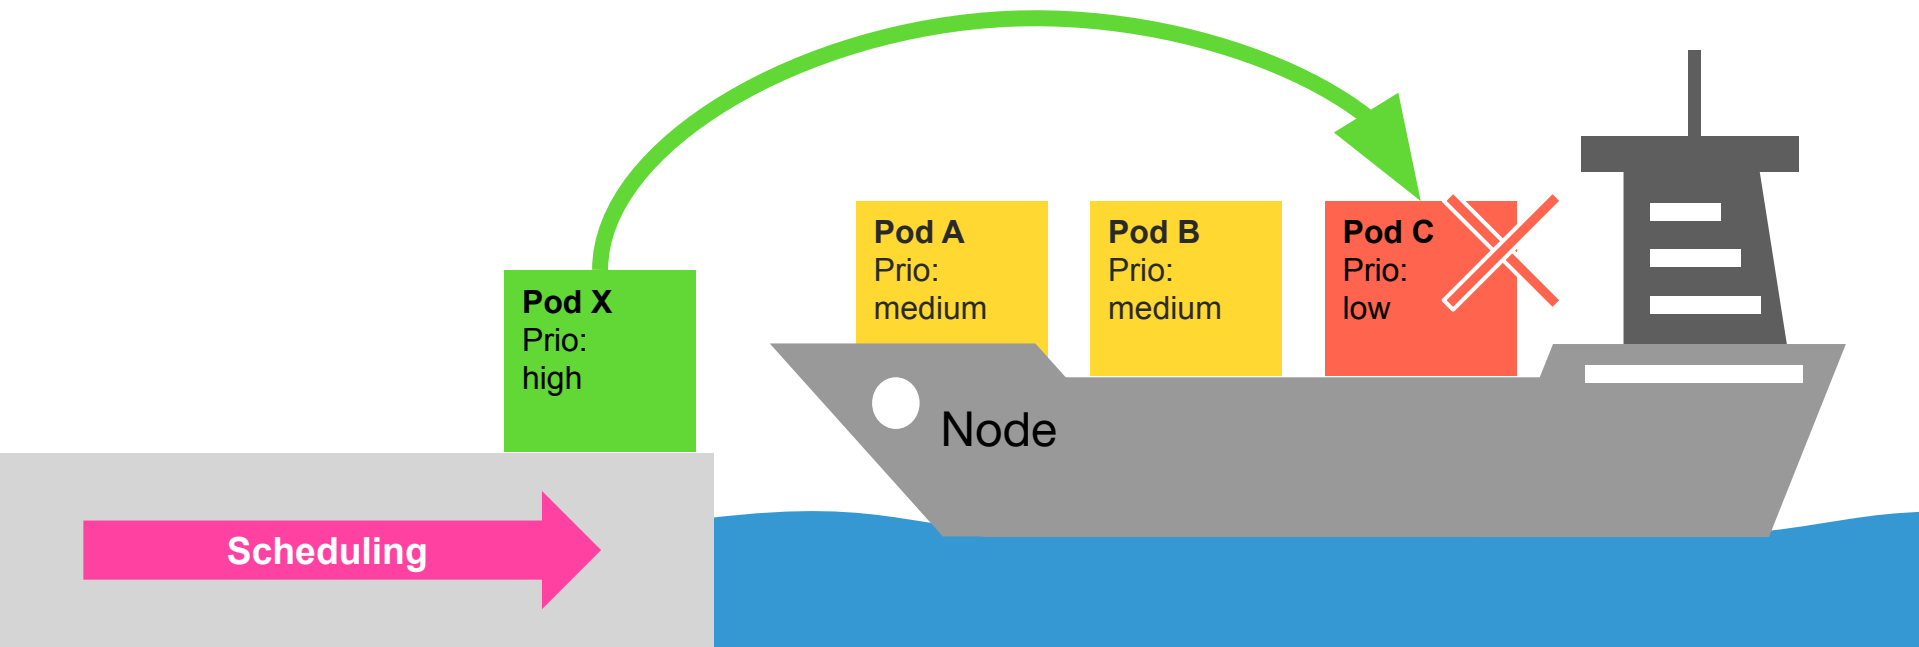
```
apiVersion: scheduling.k8s.io/v1  
kind: PriorityClass  
metadata:  
  name: high  
value: 1000000  
description: "For most important apps"
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  containers:  
    - name: nginx  
      image: nginx  
  priorityClassName: high
```




## **preemptionPolicy** - Eviction of pods with lower priority

- protects high prio Pods from being evicted



# Probes



## startupProbes

- at start time
- checks if application is started properly
- executed before any other probe
- deactivates other probes if unsuccessful

## readinessProbes

- at start time
- Pod “ready” when probe successful
- receives Traffic when “ready”


## livenessProbes

- regular checks at runtime
- if probe not successful, Pod is restarted

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/115-probes>

# Horizontal Pod Autoscaler



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/120a-hpa>



# Pod Disruption Budget



# Hands-on


<https://github.com/syseleven/containerday-workshops/tree/main/120b-pdb>

# Sidecar Containers


# Authentication and TLS offloading



# Backup




# Metrics & Healthchecks



# Logs

# Repair & Maintenance





# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/200-sidecar>

# Init Containers

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/300-init-container>

# Security for Pods & Containers

# Security Context

## Pod level

- fsGroup
- fsGroupChangePolicy
- runAsGroup
- runAsNonRoot
- runAsUser
- seLinuxOptions
- seccompProfile
- supplementalGroups
- sysctls

## Container level

- allowPrivilegeEscalation
  - capabilities
  - privileged
  - procMount
  - readOnlyRootFilesystem
  - runAsGroup
  - runAsNonRoot
  - runAsUser
  - seLinuxOptions
  - seccompProfile
- more specific than Pod level



# Security Context

**control user permissions of the running process**

- `runAs...`

**use security frameworks & kernel features**

- `secompProfile`
- `seLinux`
- `allowPrivilegeEscalation`
- `capabilities`

# Pod Security Standards (PSS)

- **successor to Pod Security Policies (PSP)** (removed in v1.25)
- works on **cluster level** or in **single namespace**
- can be added and removed at any time

# Pod Security Standards (PSS)

- 3 profiles:
  - **privileged** (unrestricted)
  - **baseline** (minimal security)
  - **restricted** (heavily restricted, best-practice security)



# Pod Security Standards (PSS)



- 3 modes:
  - **warn** (Pod will be created, user visible warning)
  - **audit** (Pod creation permitted but will be logged)
  - **enforce** (Pod cannot be created)

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/400-security>

# Sandboxing

# Sandboxing



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/410-sandboxing>

# Persistence

# PersistentVolumeClaim



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/350-persistent-storage>



# 3rd-party tools



# External-DNS




- synchronize DNS records
- multi-provider capable
- controlled by K8s resources (svc, ing)
- annotations
- automation

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/500-external-dns>

## Recap: Without ingress controller



# Ingress Controller

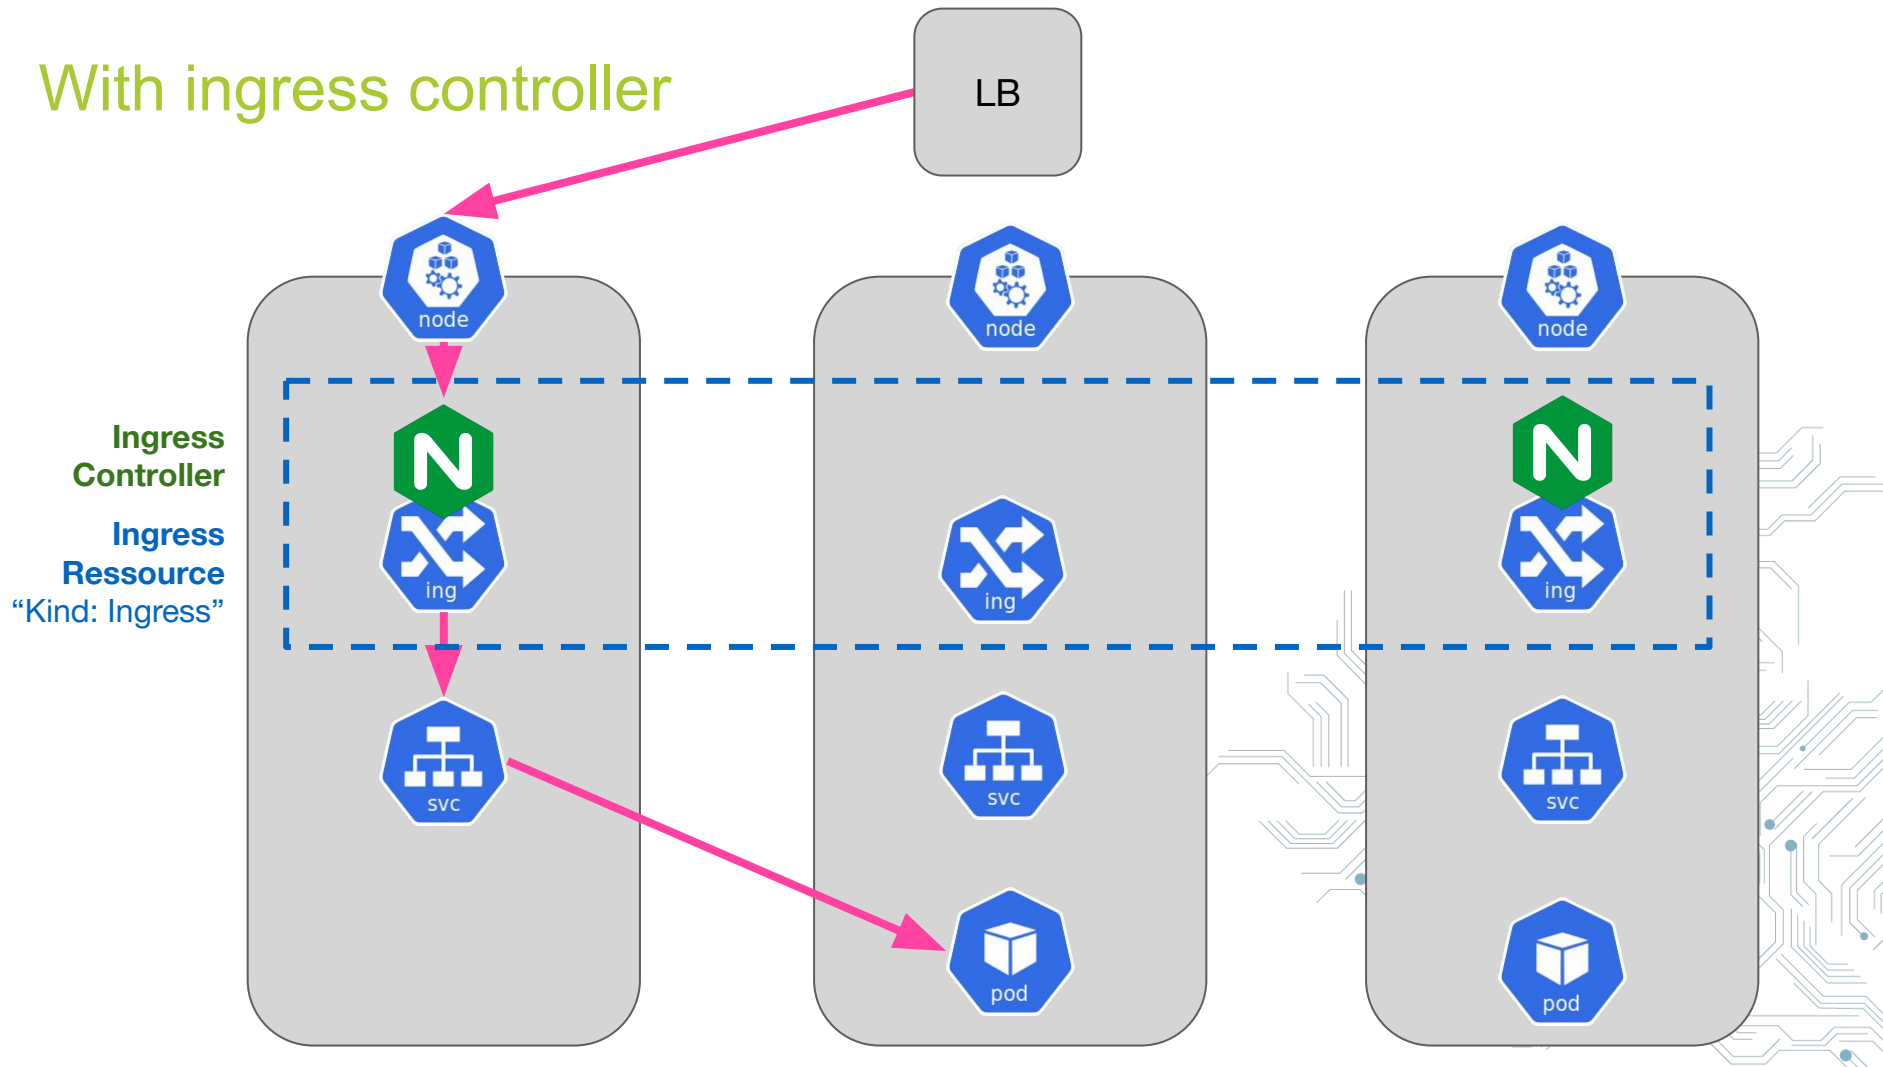


# ingress-nginx



- works on Layer7 HTTP/S
- routing based on
  - domain name
  - URL path
- SSL offloading
- basic authentication
- IP filtering

## With ingress controller





# Ingress Controller vs. Ingress Resource

**Ingress  
Controller**

**Ingress  
Resource**  
“Kind: Ingress”



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/600-ingress-nginx>

# Cert-Manager



- Let's Encrypt
- HTTP / DNS challenge
- wildcard certificates (DNS challenge)
- auto renew
- controlled by K8s resource (ing)

# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/700-cert-manager>

external-dns, ingress-nginx  
& cert-manager  
in action

1. create ingress resource



2. create DNS record



3. request certificate




# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/800-deployment-with-ingress-and-tls>



canary deployments  
in action

# Canary deployments



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/810-canary-deployments>

one **Helmfile** to rule them all

# Tooling: Helm

## Anatomy of a Helm Chart

### Chart.yaml

- Release Name
- Chart Version
- App Version

### values.yaml

- default config parameters

Deployment

Secret

Config  
Map

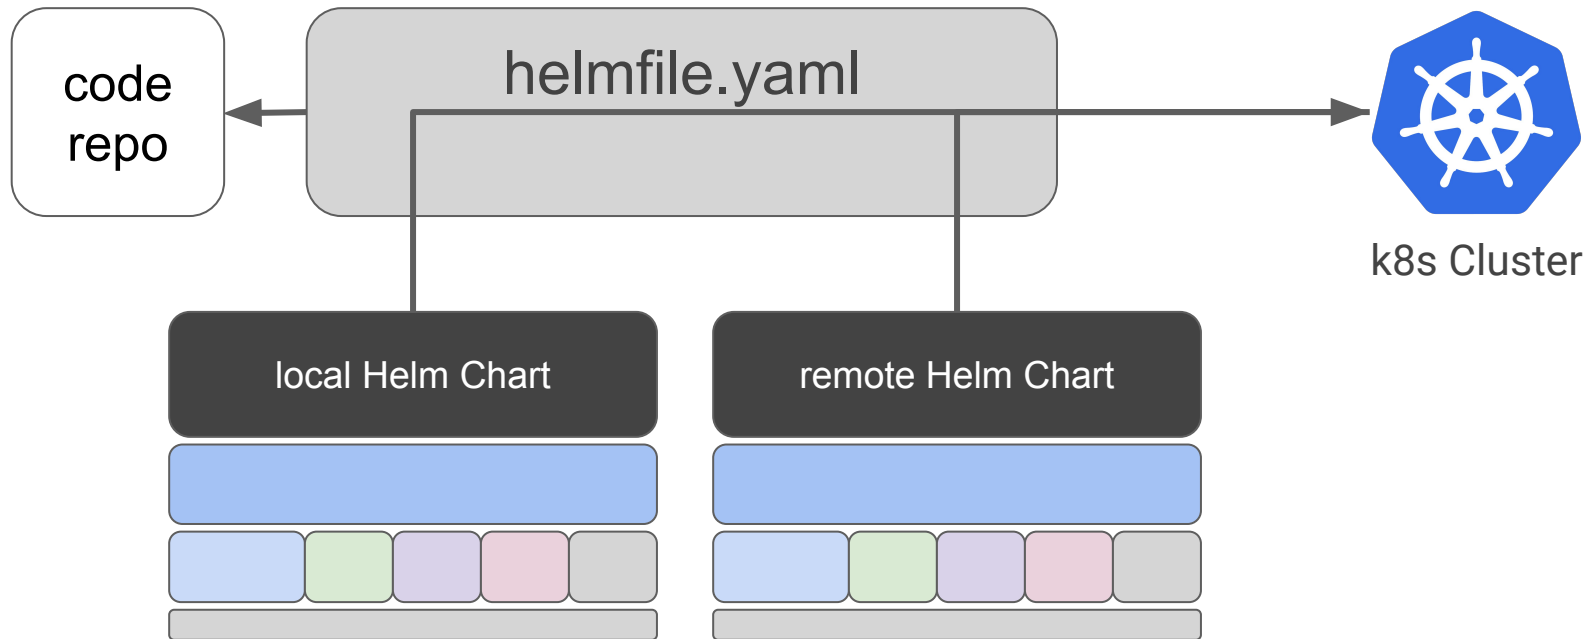
Volume

Ingress

... other templates



# Tooling: Helmfile



# Hands-on

<https://github.com/syseleven/containerday-workshops/tree/main/900-helmfile>

# reactive & supportive measures



Observability, Monitoring & Alerting



Backup & Recovery



# laaC, automation, continuous deployment



Questions?

Thank you!

[sys eleven.de/trial-metakube](https://sys eleven.de/trial-metakube)